



# Class Announcements

## 1) Install Docker

- <https://umd.instructure.com/courses/1358712/assignments>

| ▼ Assignments |  |                      | + | ⋮ |
|---------------|--|----------------------|---|---|
| ⋮             |  Set up Linux env on your laptop for the projects | Due Feb 9 at 11:59pm | ✓ | ⋮ |
| ⋮             |  Git and GitHub tutorials                         | Due Feb 9 at 11:59pm | ✓ | ⋮ |

## 2) Do the Git and GitHub tutorial

## 3) Quiz about Git next week

## 4) Midterm class project


- Sign up to Sorrentum on GitHub
- Study and experiment with different big data technology
- We will give you different examples
  - Individual project
  - Different complexity

# **UMD DATA605 - Big Data Systems**

## **DevOps with Docker**

# Docker - Resources

- We will use Docker during the class project and most tutorials
- Concepts in the slides
- Class tutorials:
  - [tutorial\\_docker](#)
  - [tutorial\\_docker\\_compose](#)
- Web resources:
  - [Docker Tutorial for beginners](#)
  - <https://labs.play-with-docker.com/>
  - <https://training.play-with-docker.com>
  - [A Beginner-Friendly Introduction to Containers, VMs and Docker](#)
  - [Official Docker Getting Started Tutorial](#)
- Mastery:
  - Poulton, [Docker Deep Dive: Zero to Docker in a single book](#), 2020




# Application Deployment

- **For (almost all) Internet companies, the application *is* the business**
  - If the application breaks, the business stops working
  - E.g., Amazon, Google, Facebook, on-line banks, travel sites (e.g., Expedia), ..., OpenAI
- **Problem**
  - How to release / deploy / manage / monitor applications?
- **Solutions**
  - Before 2000s: “bare-metal era”
  - 2000s-2010s: “virtual machine era”
  - > ~2013: “container era”

# DevOps

- **DevOps** = set of practices that combines:
  - Software development (*dev*)
  - IT operations (*ops*)
- **Containers revolutionized DevOps**
  - Enable true independence between application development and IT ops
    - One team creates an application
    - Another team deploys and manages the applications
  - Create a model for better collaboration (fewer conflicts) and innovation
    - IT: “It doesn’t work!”
    - Devs: “What? It works for me”



- Plan
- Code
- Build
- Test
- Release
- Deploy
- Operate
- Monitor

# Run on bare metal

- **< 2000s**
  - Running one or few applications on each server (without virtualization)
- **Pros**
  - No virtualization overhead
- **Cons**
  - Not safe / not secure since no separation between applications
  - Expensive
- **Expensive / low efficiency**
  - IT would buy a new server for each application
  - Difficult to spec out the machine → buy “big and fast servers”
    - Overpowered servers operating at 5-10% of capacity
    - Tons of money in the 2000 DotCom boom was spent on machines and networks
  - It kind of came back in 2020 but with different use cases in Cloud Computing
- **Winners:** Cisco, Sun, Microsoft




# Virtual Machine Era

- **Circa 2000-2010: Virtual Machine**
  - Virtual machine technology = run multiple copies of OSes on the same hardware
- **Pros**
  - VM runs safely and securely multiple applications on a single server
  - IT could run apps on existing servers with spare capacity
- **Cons**
  - Every VM requires an OS (waste of CPU, RAM, and disk)
  - Buy an OS license
  - Monitor and patch each OS
  - VMs are slow to boot
- **Winners:** VMWare, RedHat, Citrix



# Containers Era

- **Circa 2013: Docker becomes ubiquitous**
- **Docker**
  - Didn't invent containers
  - Made containers simple and mainstream
- Linux supported containers for some time
  - Kernel namespaces
  - Control groups
  - Union filesystems
- **Pros**
  - Containers are fast and portable
  - Containers don't require full-blown OS
  - All containers run on a single host
  - Reduce OS licencing cost
  - Reduce overhead of OS patching and maintenance
- **Cons**
  - CPU overhead
  - Toolchain to learn / use
- **Winners:** AWS, Microsoft Azure, Google (not Docker Inc.)





# Serverless Computing

- **Containers run in an OS, OS runs on a host**

- **Where is the host running?**

- Local (your laptop)
    - On premise (your own computer in a rack)
    - Cloud instance (e.g., AWS EC2 instance)

- **How is the host running?**

- On bare-metal server
    - On a virtual machine
    - On a virtual machine running a virtual machine

- **Serverless computing**

- As long your application runs somewhere, you don't care “how” or “where”
  - E.g., AWS Lambda

# HW vs OS Virtualization

- **Hypervisor performs HW virtualization**

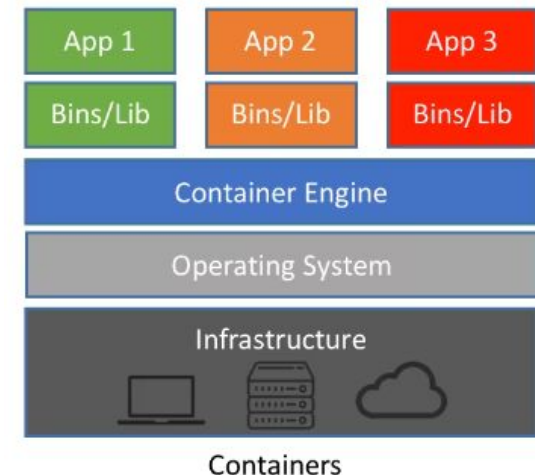
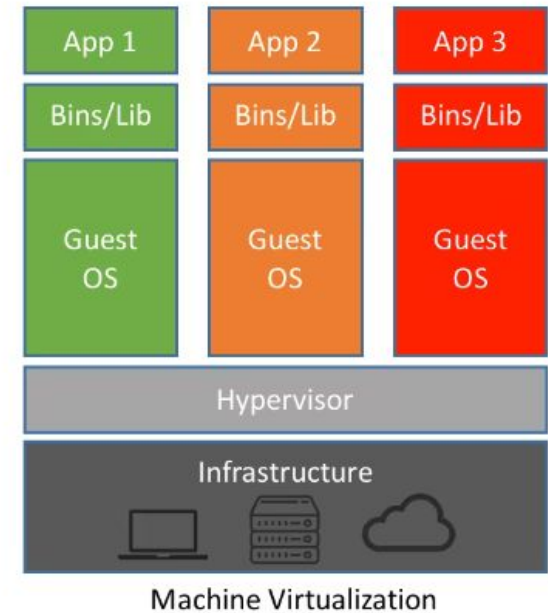
- Carves out physical hardware resources into VMs
- Resources (CPUs, RAM, storage) are allocated to a VM
- It's like having multiple computers

- “Virtual machine tax”

- To run 3 apps, you need 3 VMs
  - Each VM requires time to start
  - Consumes CPU, RAM, storage
  - VM needs a OS license
  - VM / OS need admins, patching
- You just want to run 3 apps!


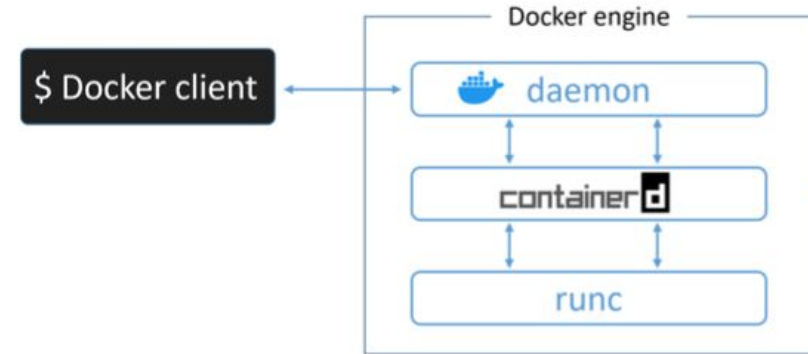
- **Containers perform OS virtualization**

- It's like having multiple OSes



# Docker: Client-Server

- Client-Server architecture
- **Docker client**
  - Command line interface
  - Communicate with the server through IPC socket
    - E.g., `/var/run/docker.sock` or IP port
- **Docker engine**
  - Run and manage containers
  - Modular and built from several OCI-compliant sub-systems
    - E.g., Docker daemon, `containerd`, `runc`, plug-ins for networking and storage



# Docker Architecture

- **Docker run-time**

- **runc**: start and stop containers
- **containerd**
  - Pull images
  - Create containers, volumes, network interfaces

- **Docker engine**

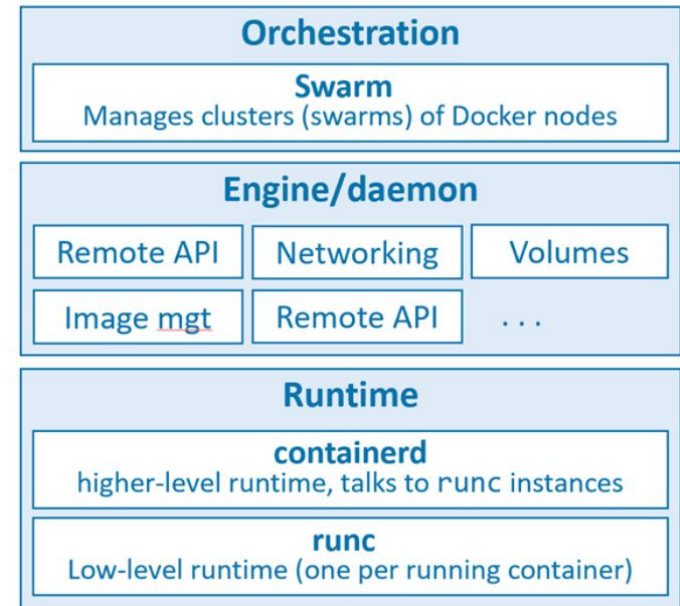
- **dockerd**
  - Expose remote API
  - Manage images, volumes, networks

- **Docker orchestration**

- **docker swarm**
- Manage clusters of nodes
- Replaced by Kubernetes

- **Open Container Initiative (OCI)**

- Standardize low-level components of container infrastructure
- E.g., image format, run-time API
- “Death” of Docker



# Docker Container

- **Docker Container**

- Unit of computation
- Lightweight, stand-alone, executable software package
- Include everything needed to run
  - E.g., code, runtime / system libraries, settings
- It is a run-time object
  - vs Docker images are built-time objects
  - Like program running (container) vs program code (image)




# Docker Image

- **Docker Image**

- Unit of deployment
- Contain everything needed by an app to run
  - Application code
  - Application dependencies
  - Minimal OS support
  - File system
- Users can
  - Build images from Dockerfiles
  - Pull pre-built images from a registry
- Multiple layers stacked on top of each other
  - Typically few 100s MBs

# Docker Image Layers

- **Docker image** is a configuration file that lists the layers and some metadata
  - It is composed of read-only layers
  - Each layer is independent from each other
  - Each layer comprises of many files
- **Docker driver**
  - Stacks these layers as a unified filesystem
  - Implements a copy-on-write behavior
  - Files from the top layers can obscure the files from the bottom layers
- **Layer hash**
  - Each layer has an hash based on its content
  - Layers are pulled and pushed compressed
- **Image hash**
  - Each image has an hash
  - The hash is function of the config file and of the layers
  - When an image changes, a new hash is generated



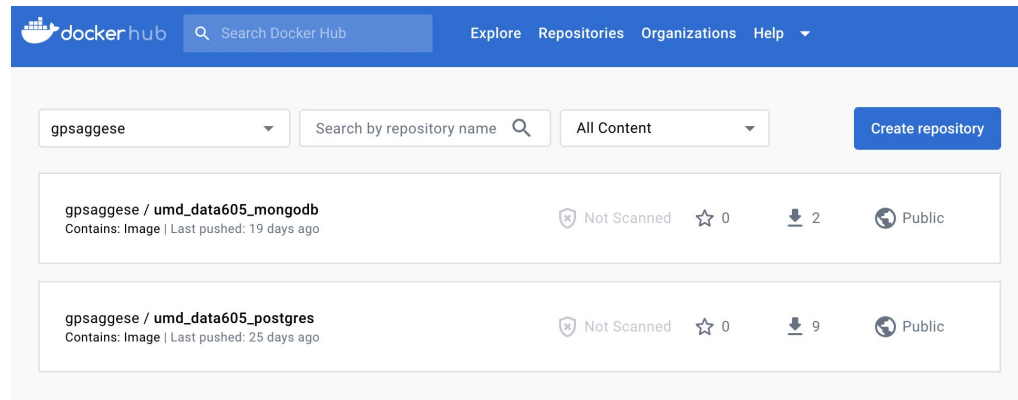

# Docker: Container Data

- A container has access to different data
- **Container storage**
  - It is a copy-on-write layer in the image
  - It is ephemeral (only temporary data)
  - Data inside of containers is persisted as long as the container is not killed
    - If you stop or pause a container data is not lost
  - Containers are designed to be immutable
    - It's not good practice to write “persistent” data into containers
- **Bind-mount a local dir**
  - = a local dir is mounted to a dir inside a container
- **Docker volumes**
  - Docker provides volumes that exist separately from the container
    - E.g., to store the content of a Postgres DB
  - State is permanent across container invocations
  - Can be shared across containers




# Docker Repos

- **Docker Repo (Registry)**
  - Store Docker images
    - `<registry>/<repo>:<tag>`
    - `docker.io/alpine:latest`
    - E.g., DockerHub, AWS ECR
  - Some repos are vetted by Docker
  - Unofficial repos shouldn't be trusted
  - E.g., <https://hub.docker.com/>



# Devops = Devs + Ops



- **Devs**

- Implement the app
  - E.g., Python, virtual env
- Containerize the app
  - Create Dockerfile
  - Contain the instruction on how to build an image
- Build image
- Run the app as a container
- Test “locally”

- **Ops**

- Download container images
  - Contain filesystem, application, app dependencies
- Start / destroy containers
- In case of issues, it's easy to reproto the problem
  - “Here is the log”
  - Run command line
  - Deploy on a test system and debug

# Containerizing an App

- = create a container with your app inside
- Develop your application code using the needed dependencies
  - E.g., install dependencies
    - Directly inside a container
    - Inside a virtual env
- Create a Dockerfile describing:
  - your app
  - its dependencies
  - how to run it
- Build image with **docker image build**
- (Optional) Push image to a Docker image registry
- Run / test container from image
- Distribute your app as a container (no installation)

# Building a Container

- **Dockerfile**

- Describe how to create a container

- **Build context**

- `> docker build -t web:latest .`
- `.` is the build context
- Directory containing the application and what's needed to build it
- Sent to Docker engine to build the application
- Typically the Dockerfile is in the root directory of the build context

# Dockerfile Example

```
FROM python:3.8-slim-buster
```

```
LABEL maintainer="gsaggese@umd.edu"
```

```
WORKDIR /app
```

```
COPY requirements.txt requirements.txt
```

```
RUN pip3 install -r requirements.txt
```

```
COPY . .
```

```
CMD ["python3", "-m", "flask", "run", "--host=0.0.0.0"]
```

# Docker Tutorial

- [tutorial\\_docker.md](#)

# Docker Compose

- **Manage multi-container apps running on a single node**
  - Describe app in a single *declarative* configuration YAML file
    - Instead of scripts with long Docker commands
  - Compose talks to Docker API to achieve what you requested
  - E.g., you need a client app and Postgres DB
  - E.g., microservices
    - Web front-end
    - Ordering
    - Back-end DB
- In 2020 Docker Compose has become an open standard for “code-to-cloud” process
- **Manage multi-container apps running on multiple hosts**
  - Docker Stacks / Swarm
  - Kubernetes

# Docker Compose: Commands

> docker compose --help

Usage: docker compose [OPTIONS] COMMAND

Options:

|                           |  |
|---------------------------|--|
| --env-file string         | Specify an alternate environment file. |
| -f, --file stringArray    | Compose configuration files            |
| -p, --project-name string | Project name                           |

Commands:

|         |   |
|---------|---|
| build   | Build or rebuild services   |
| convert | Converts the compose file to platform's canonical format                |
| cp      | Copy files/folders between a service container and the local filesystem |
| create  | Creates containers for a service.                                       |
| down    | Stop and remove containers, networks                                    |
| events  | Receive real time events from containers.                               |
| exec    | Execute a command in a running container.                               |
| images  | List images used by the created containers                              |
| kill    | Force stop service containers.  |
| logs    | View output from containers   |
| ls      | List running compose projects   |
| pause   | Pause services  |
| port    | Print the public port for a port binding.                               |
| ps      | List containers   |
| pull    | Pull service images   |
| push    | Push service images   |
| restart | Restart containers  |
| rm      | Removes stopped service containers                                      |
| run     | Run a one-off command on a service.                                     |
| start   | Start services  |
| stop    | Stop services   |
| top     | Display the running processes   |
| unpause | Unpause services  |
| up      | Create and start containers   |
| version | Show the Docker Compose version information                             |



# Docker Compose: Tutorial Example

- The default name for a Compose file is ``docker-compose.yml``
  - You can specify ``-f`` for custom filenames
- **Top-level keys** are:
  - ``version``:
    - Mandatory first line to specify API version
    - Ideally always use the latest version
    - Typically 3 or higher
  - ``services``:
    - Define the different microservices
  - ``networks``:
    - Creates new networks
    - By default it creates a ``bridge`` network to connect multiple containers on the same Docker host
  - ``volumes``:
    - Creates new volumes
- **Key in services** describe a different “service” in terms of container
  - **Inner keys** specify the params of Docker run command

```
version: "3.8"
```

```
services:
```

```
  web-fe:
```

```
    build: .
```

```
    command: python app.py
```

```
    ports:
```

```
      - target: 5000
```

```
        published: 5001
```

```
    networks:
```

```
      - counter-net
```

```
    volumes:
```

```
      - type: volume
```

```
        source: counter-vol
```

```
        target: /code
```

```
  redis:
```

```
    image: "redis:alpine"
```

```
    networks:
```

```
      counter-net:
```

```
networks:
```

```
  counter-net:
```

```
volumes:
```

```
  counter-vol:
```

# Docker Compose: Tutorial

- Example taken from <https://github.com/nigelpoulton/counter-app>
- [tutorial\\_docker\\_compose](#)
  - > `cd tutorials/tutorial_docker_compose`
  - > `vi tutorial_docker_compose.mdx`

2023 Version

# Class Announcements

## 1) Project teams

- Posted [UMD DATA605 - Class Project Teams - Spring 2023](#)
- If your name is not on any of the teams, please send me an email
- No midterm or final exam, complete class project to get a grade

## 2) Team composition

- Teams based on your self-assessed skills
- In each group there should be someone with experience with Git, Docker, Python, and so on
- Teams are not perfect, but none of your team at your future jobs will be
- Working in a team is a skill that takes a long time to hone: let's start practicing it

## 3) Class project complexity

- Projects have about the same complexity
- If not, we will try to account for this when grading
- Projects were assigned randomly to the teams

# Class Announcements

## 4) Add personal info

- Fill the spreadsheet [UMD DATA605 - Class Project Teams - Spring 2023](#) with your information
- Email
- GitHub username (free account at <https://github.com>)
- Telegram Handle for IM (free account at <https://telegram.org>)

## 5) Next steps

- Read carefully [UMD DATA605 - Class Project](#)
- If things are not clear send us an email or add a comment to the Google Doc
- Read the Google doc corresponding to your assigned project from [UMD DATA605 - Class Project Teams - Spring 2023](#)
- Read the first Deliverable 1 from [UMD DATA605 - Class Project](#)

## 6) Golden rule

- *Treat others how you want to be treated*
  - Everybody comes from a different place and different skill level, somebody has a job, somebody has a full-time work
- *If you want to go fast, go alone; if you want to go far, go together*