

UMD DATA605 - Big Data Systems

NoSQL Document Stores

MongoDB

CouchDB

Dr. GP Saggese

gsaggese@umd.edu

with thanks to Profs

Alan Sussman (UMD)

Amol Deshpande (UMD)

Nguyen Vo (Utah State U.)

Kathleen Durant (Northeastern U.)

UMD DATA605 - Big Data Systems

NoSQL Document Stores

MongoDB


Couchbase

Key-Value Store vs Document DBs

- **Key-value stores**

- Basically a map or a dictionary
 - E.g., HBase, Redis
- Typically only look up values by key
 - Sometimes can do search in value field with a pattern
- Uninterpreted value (e.g., binary blob) associated with a key
- Typically one namespace for all key-values


Key-Value



- **Document DBs**

- Collect sets of key-value pairs into *documents*
 - E.g., MongoDB, CouchDB
- Documents represented in JSON, XML, or BSON (binary JSON)
- Documents organized into *collections*
 - Similar to tables in relational DBs
 - Large collections can be partitioned and indexed

Document



UMD DATA605 - Big Data Systems


NoSQL Document Stores

MongoDB

Couchbase

Resources

- All concepts in slides
- [MongoDB tutorial](#)
- Web
 - <https://www.mongodb.com/>
 - [Official docs](#)
 - [pymongo](#)
- Book
 - [Seven Databases in Seven Weeks, 2e](#)



MongoDB



- Developed by MongoDB Inc.
 - Founded in 2007
 - Based on DoubleClick experience with large-scale data
 - Mongo comes from “hu-mongo-us”
- One of the most used NoSQL DBs (if not the most used)
- **Document-oriented NoSQL database**
 - Schema-less
 - No Data Definition Language (DDL), like for SQL
 - You can store maps with any keys and values
 - Application tracks the schema, mapping between documents and their meaning
 - Keys are hashes stored as strings
 - Document Identifiers `_id` created for each document (field name reserved by Mongo)
 - Values use BSON format
 - Based on JSON (B stands for Binary)
- Written in C++
- Supports APIs (drivers) in many languages
 - E.g., JavaScript, Python, Ruby, Java, Scala, C++, ...

MongoDB: Example of Document

- **A document is a JSON data structure**
- It corresponds to a row in a relational DB
 - Without schema
 - Primary key is id
 - Values nested to an arbitrary depth

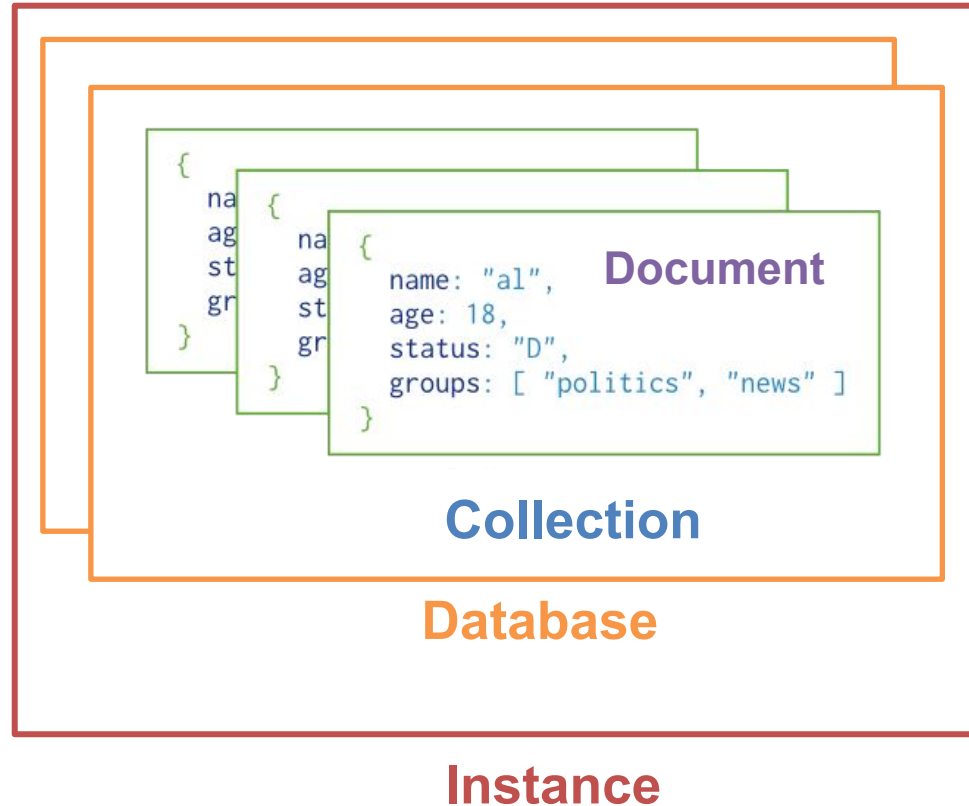
```
{
  "_id" : ObjectId("4d0b6da3bb30773266f39fea"),
  "country" : {
    "$ref" : "countries",
    "$id" : ObjectId("4d0e6074deb8995216a8309e")
  },
  "famous_for" : [
    "beer",
    "food"
  ],
  "last_census" : "Sun Jan 07 2018 00:00:00 GMT -0700 (PDT)",
  "mayor" : {
    "name" : "Ted Wheeler",
    "party" : "D"
  },
  "name" : "Portland",
  "population" : 582000,
  "state" : "OR"
}
```

MongoDB: Functionalities

- **Design goals**
 - Performance
 - Availability / scalability
 - Rich data storage (not rich querying!)
- **Dynamic schema**
 - No DDL (Data Definition Language)
 - Secondary indexes
 - Query language via an API
- **Several levels of data consistency**
 - E.g., atomic writes and fully-consistent reads (at document level)
- **No joins nor transactions across multiple documents**
 - Makes distributed queries easy and fast
- **High availability through replica sets**
 - E.g., primary replication with automated failover
- **Built-in sharding**
 - Horizontal scaling via automated range-based partitioning of data
 - Reads and writes distributed over shards

MongoDB: Hierarchical Objects

- A Mongo **instance** has:
 - Zero or more “databases”
 - Mongo instance ~ Postgres instance
- A Mongo **database** has:
 - Zero or more “collections”
 - Mongo collection ~ Postgres tables
 - Mongo database ~ Postgres database
- A Mongo **collection** has:
 - Zero or more “documents”
 - Mongo document ~ Postgres rows
- A Mongo **document** has:
 - One or more “fields”
 - It has always primary key `_id`
 - Mongo field ~ Postgres columns



From <https://www.mongodb.com/docs/manual/core/data-modeling-introduction>

Relational DBs vs MongoDB: Terms and Concepts

RDBMS Concept	MongoDB Concept	Meaning in MongoDB
database	database	Container for collections
relation / table / view	collection	Group of documents
row / instance	document (BSON)	Group of fields
column / attribute	field	A name-value pair
index	index	Automatic
primary keys	<u>_id</u> field	Always the primary key
foreign key	reference	Pointers
table joins	embedded documents	Nested name-value pairs

```
{
  "_id" : ObjectId("4d0b6da3bb30773266f39fea"),
  "country" : {
    "$ref" : "countries",
    "$id" : ObjectId("4d0e6074deb8995216a8309e")
  },
  "famous_for" : [
    "beer",
    "food"
  ],
  "last_census" : "Sun Jan 07 2018 00:00:00 GMT -0700 (PDT)",
  "mayor" : {
    "name" : "Ted Wheeler",
    "party" : "D"
  },
  "name" : "Portland",
  "population" : 582000,
  "state" : "OR"
}
```

Relational vs Document DB: Workflows

- **Relational DBs**


- E.g., PostgreSQL
- Know what you want to store
 - Tabular data
- Do not know how to use it
 - Static schema allows query flexibility (e.g., joins)
- Complexity is at insertion time
 - Decide how to represent the data (i.e., schema)

- **Document DBs**

- E.g., MongoDB
- No assumptions on what to store
 - E.g., irregular JSON data
- Know a bit how to access data
 - You want to access the data by key
 - E.g., it's a nested key-value map
- Complexity is at access time
 - Get the data from the server
 - Process data on the client side

Why Use MongoDB?

- Simple to query
 - Do the work on client side
- It's fast
 - 2-10x faster than Postgres
- Data model / functionalities suitable for most web applications
 - Semi-structured data
 - Quickly evolving systems
- Easy and fast integration of data
- Not well suited for heavy and complex transactions systems
 - E.g., banking system



MongoDB: Data Model

- **Documents** are composed of field and value pairs
 - **Field names** are strings
 - **Values** are any BSON type
 - Arrays of documents
 - Native data types
 - Other documents
- E.g.,
 - **_id** holds an ObjectId
 - **name** holds a document that contains the fields **first** and **last**
 - **birth** and **death** are of Date type
 - **contribs** holds an array of strings
 - **views** holds a value of the NumberLong type

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

← field: value
← field: value
← field: value
← field: value

```
{  
  _id: ObjectId("5099803df3f4948bd2f98391"),  
  name: { first: "Alan", last: "Turing" },  
  birth: new Date('Jun 23, 1912'),  
  death: new Date('Jun 07, 1954'),  
  contribs: [ "Turing machine", "Turing test", "Turingery" ],  
  views : NumberLong(1250000)  
}
```

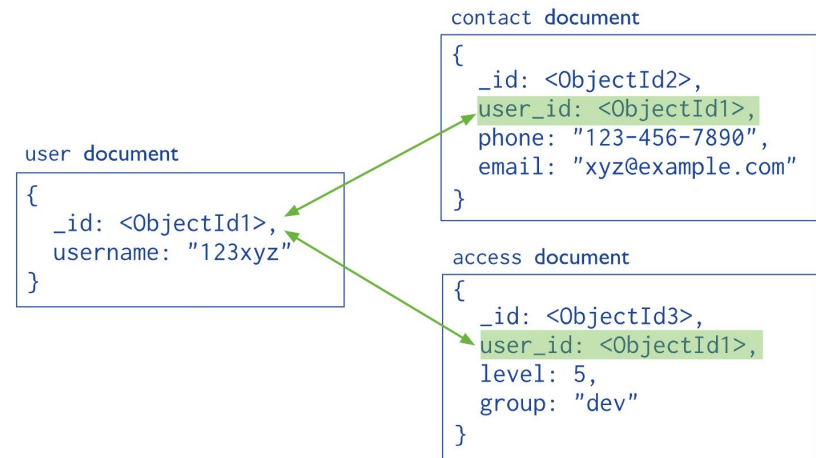
MongoDB: Data Model

- Documents can be nested
 - Embedded sub-document
- **Denormalized data models**
 - Store multiple related pieces of information in the same record
 - Conceptually is the result of a join operation
- **Normalized data models**
 - Eliminate duplication
 - Represent many-to-many relationships

```
{
  _id: <ObjectId>,
  username: "123xyz",
  contact: {
    phone: "123-456-7890",
    email: "xyz@example.com"
  },
  access: {
    level: 5,
    group: "dev"
  }
}
```

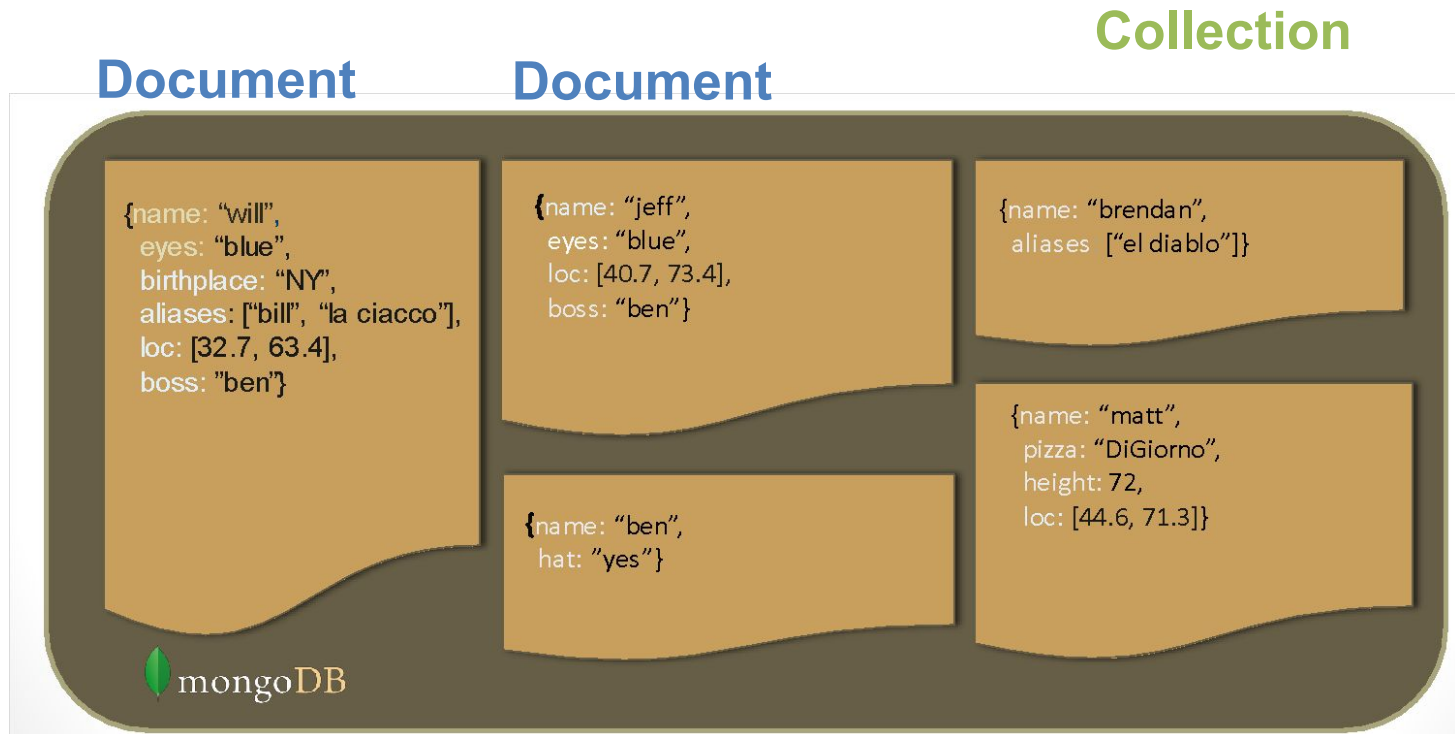
Embedded sub-document

Embedded sub-document



Schema Free

- MongoDB does not need any pre-defined data schema
- Every **document** in a **collection** can have different fields and values
 - No need for NULL values / union of fields like in relational DBs
- E.g., dishomogeneous data instances



JSON Format

- JSON = JavaScript Object Notation
- Data is stored in field / value pairs
- A field / value pair consists of:

- A field name (always a string)
- Followed by a colon :
- Followed by a typed value

`"name": "R2-D2"`

- Data in documents is separated by commas ,

`"name": "R2-D2", race : "Droid"`

- Curly braces {} hold documents

`{"name": "R2-D2", race : "Droid", affiliation: "rebels"}`

- An array is stored in brackets []

`[{"name": "R2-D2", race: "Droid", affiliation: "rebels"},
{"name": "Yoda", affiliation: "rebels"}]`

- Supports:


- Embedding of nested objects within other objects
- Just references

BSON Format

- Binary-encoded serialization of JSON-like documents
 - <https://bsonspec.org>
- Zero or more key/value pairs are stored as a single entity
 - Each entry consists of:
 - a field name (string)
 - a data type
 - a value
- Similar to protocol buffer, but more schema-less
- Large elements in a BSON document are prefixed with a length field to facilitate scanning
- MongoDB understands the internals of BSON objects, even nested ones
 - Can build indexes and match objects against query expressions for BSON keys

ObjectId

- Each JSON data contains an `_id` field of type ObjectId
 - Same as a `SERIAL` constraint incrementing a numeric primary key in PostgreSQL
- An ObjectId is 12 bytes, composed of:
 - a timestamp
 - client machine ID
 - client process ID
 - a 3-byte auto-incremented counter
- Each Mongo process can handle its own ID generation without colliding
 - Mongo has a distributed nature
- Details [here](#)



Indexes

- **Primary index**
 - Automatically created on the `_id` field
 - B+ tree indexes
- **Secondary index**
- Users can create secondary indexes to:
 - Improve query performance
 - Enforce unique values for a particular field
- Single field index and compound index (like SQL)
 - Order of the fields in a compound index matters
- Sparse property of an index
 - The index contains only entries for documents that have the indexed field
 - Ignore records that do not have the field defined
- Reject records with duplicate key value if an index is unique and sparse
- Details at <https://www.mongodb.com/docs/manual/indexes/>

CRUD Operations

- CRUD = Create, Read, Update, Delete
- **Create**
 - `db.collection.insert(<document>)`
 - `db.collection.update(<query>, <update>, {upsert: true})`
 - Upsert = update (if exists) or insert (if it doesn't)
- **Read**
 - `db.collection.find(<query>, <projection>)`
 - `db.collection.findOne(<query>, <projection>)`
- **Update**
 - `db.collection.update(<query>, <update>, <options>)`
- **Delete**
 - `db.collection.remove(<query>, <justOne>)`

Details at <https://www.mongodb.com/docs/manual/crud/>

Create Operations

- **db.collection** specifies the collection (like an SQL table) to store the document

db.collection.insert(<document>)

- Without **_id** field, MongoDB generates a unique key

db.parts.insert({type: "screwdriver", quantity: 15})

- Use **_id** field if it has a special meaning

db.parts.insert({_id: 10, type: "hammer", quantity: 1})

- Update 1 or more records in a collection satisfying **query**

db.collection.update(<query>, <update>, {upsert: true})

- Update an existing record or create a new record

db.collection.save(<document>)

- A more modern OOP-like syntax than the COBOL / FORTRAN-inspired SQL

Read Operations

- **find** provides functionality similar to **SQL SELECT** command
db.collection.find(<query>, <projection>).cursor
with:
 - <query> = where condition
 - <projection> = fields in result set
- **db.parts.find({parts: “hammer”}).limit(5)**
 - Return cursor to handle a result set
 - Can modify the query to impose limits, skips, and sort orders
 - Can specify to return the ‘top’ number of records from the result set
- **db.collection.findOne(<query>, <projection>)**

More Query Examples

SQL

```
SELECT * FROM users WHERE age>33
```

```
SELECT * FROM users WHERE age!=33
```

```
SELECT * FROM users WHERE name LIKE  
"%Joe%"
```

```
SELECT * FROM users WHERE a=1 and b='q'
```

```
SELECT * FROM users WHERE a=1 or b=2
```

```
SELECT * FROM foo  
WHERE name='bob' and (a=1 or b=2 )
```

```
SELECT * FROM users  
WHERE age>33 AND age<=40
```

Mongo

```
db.users.find({age: {$gt: 33}})
```

```
db.users.find({age: {$ne: 33}})
```

```
db.users.find({name: /Joe/})
```

```
db.users.find({a: 1, b: 'q'})
```

```
db.users.find({$or: [{a: 1}, {b: 2}]})
```

```
db.foo.find({name: "bob",  
$or: [{a: 1}, {b: 2}]})
```

```
db.users.find({'age':  
{$gt: 33, $lte: 40}})
```

- Mongo has a functional programming flavor
 - E.g., composing operators, like **\$or**

Query Operators

Command	Description
\$regex	Match by any PCRE-compliant regular expression string (or just use the // delimiters as shown earlier)
\$ne	Not equal to
\$lt	Less than
\$lte	Less than or equal to
\$gt	Greater than
\$gte	Greater than or equal to
\$exists	Check for the existence of a field
\$all	Match all elements in an array
\$in	Match any elements in an array
\$nin	Does not match any elements in an array
\$elemMatch	Match all fields in an array of nested documents
\$or	or
\$nor	Not or
\$size	Match array of given size
\$mod	Modulus
\$type	Match if field is a given datatype
\$not	Negate the given operator check

Update Operations

- `db.collection.insert(<document>)`
 - Omit the `_id` field to have MongoDB generate a unique key

```
db.parts.insert({type: "screwdriver", quantity: 15})  
db.parts.insert({_id: 10, type: "hammer", quantity: 1})
```
- `db.collection.save(<document>)`
 - Updates an existing record or creates a new record
- `db.collection.update(<query>, <update>, {upsert: true})`
 - Will update 1 or more records in a collection satisfying query
- `db.collection.findAndModify(<query>, <sort>, <update>, <new>, <fields>, <upsert>)`
 - Modify existing record(s)
 - Retrieve old or new version of the record

Delete Operations

- `db.collection.remove(<query>, <justone>)`
 - Delete all records from a collection or matching a criterion
 - `<justone>` specifies to delete only 1 record matching the criterion
- Remove all records in `parts` with `type` starting with h
 - `db.parts.remove(type: /^h/ }`
- Delete all documents in the `parts` collection
 - `db.parts.remove()`

MongoDB Features

- Document-oriented NoSQL store
- Rich querying
 - Full index support (primary and secondary)
- Fast in-place updates
- Agile and scalable
 - Replication and high availability
 - Auto-sharding
 - Map-reduce functionality
- Scale horizontally over commodity hardware
 - Horizontally = add more machines
 - Commodity hardware = relatively inexpensive servers

MongoDB vs Relational DBs

- Keep the functionality that works well in RDBMSs
 - Ad-hoc queries
 - Fully featured indexes
 - Secondary indexes
- Do not offer RDBMS functionalities that don't scale up
 - Long running multi-row transactions
 - ACID consistency
 - Joins

MongoDB Tutorial

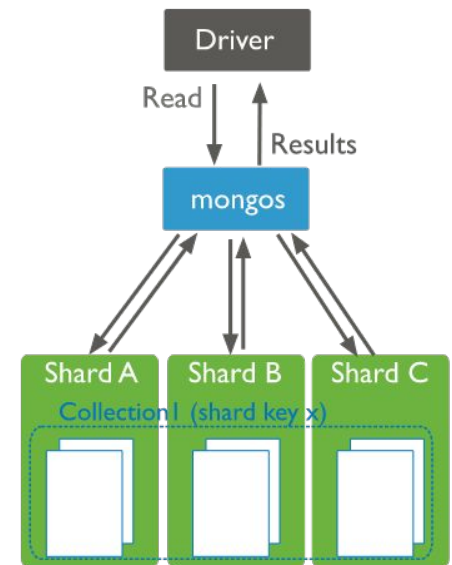
Tutorial is at [GitHub](#)

The instructions are [here](#)

```
> cd $GIT_REPO/tutorials/tutorial_mongodb  
> vi tutorial_mongo.md
```

MongoDB Processes and Configuration

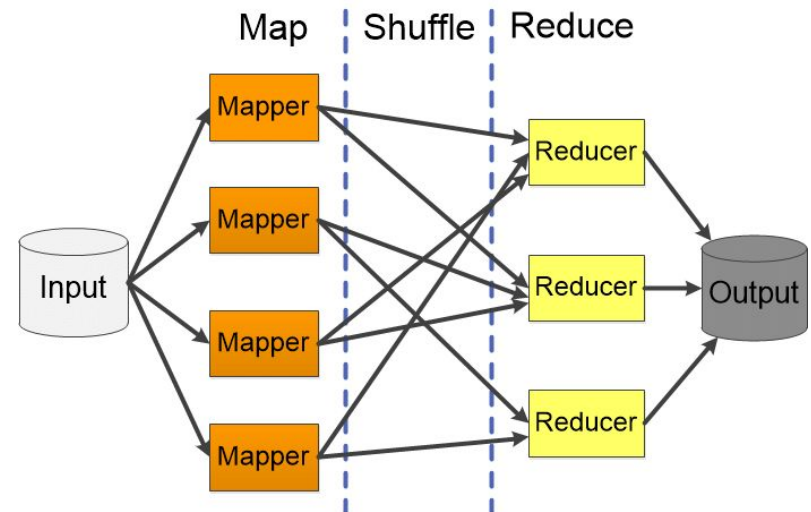
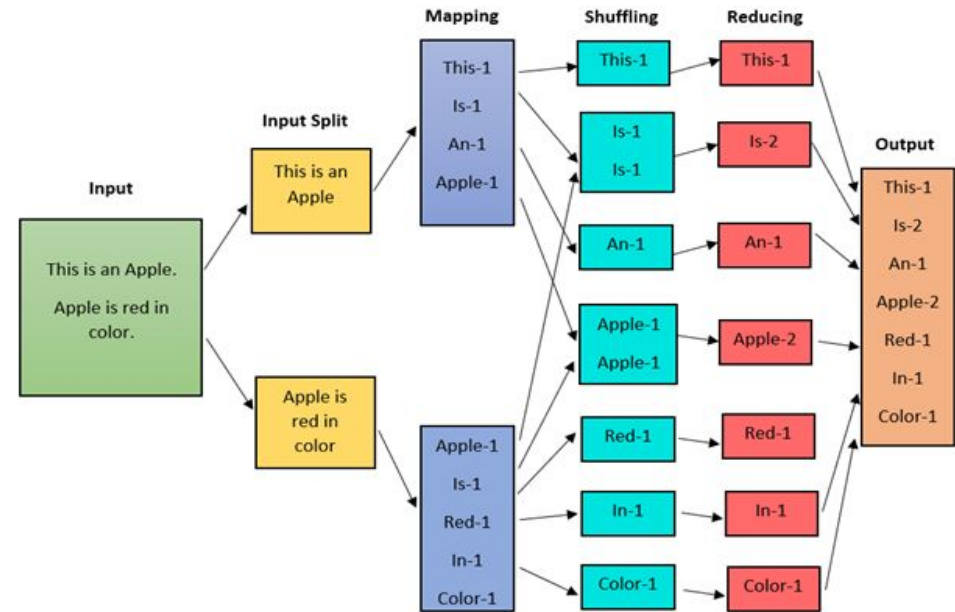
- **mongod**: database instance (i.e., a server process)
- **mongosh**: interactive shell (i.e., a client)
 - Fully functional JavaScript environment for use with a MongoDB
- **mongos**: database router
 - Process all requests
 - Decide how many and which **mongod** instances should receive the query (sharding / partitioning)
 - Collate the results
 - Send result back to the client
- You should have:
 - One **mongos** (router) for the whole system no matter how many **mongods** you have; or
 - One local **mongos** for every client if you wanted to minimize network latency



MapReduce Functionality

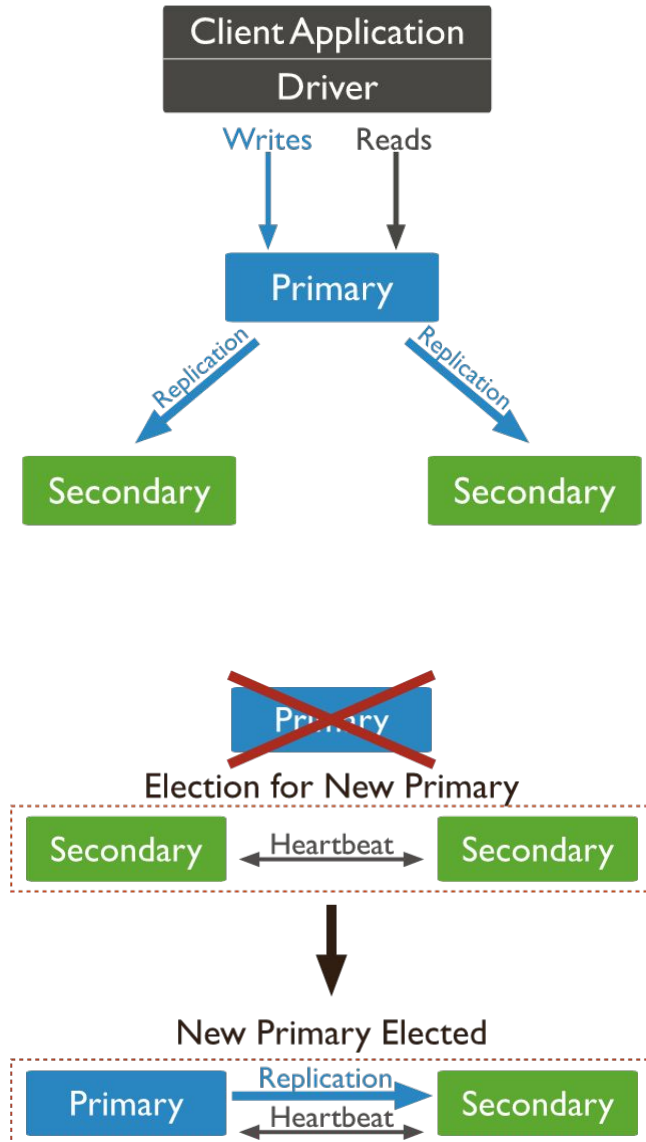
- Perform map-reduce computation given a collection of (keys, value) pairs
- Must provide at least a map function, reduction function, and the name of the result set

```
db.collection.mapReduce(
  <map_function>,
  <reduce_function>,
  {
    out: <collection>,
    query: <document>,
    sort: <document>,
    limit: <number>,
    finalize: <function>,
    scope: <document>,
    jsMode: <boolean>,
    verbose: <boolean>
  })
```



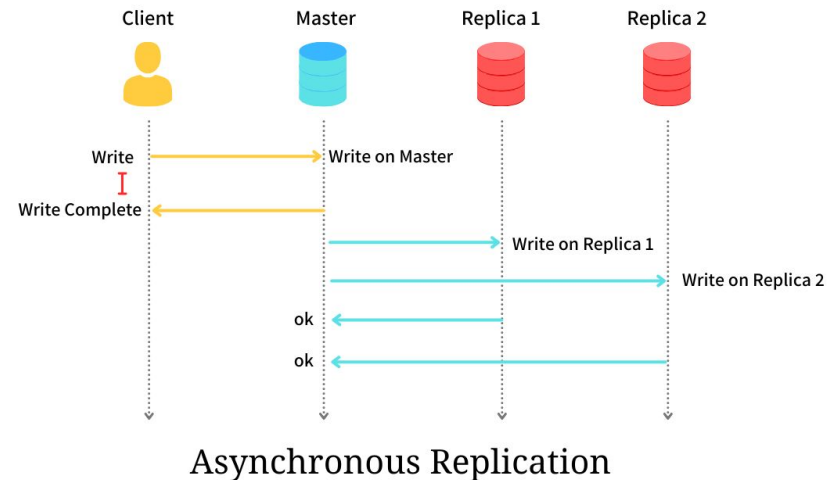
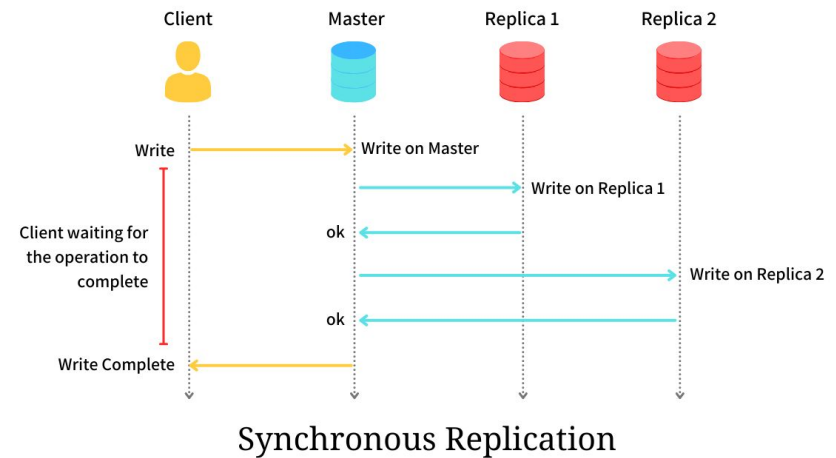
Data Replication

- **Data replication** ensure:
 - Redundancy
 - Backup
 - Automatic failover
- Replication occurs through groups of servers known as **replica sets**
 - **Primary set**: set of servers that client asks direct updates to
 - **Secondary set**: set of servers used for duplication of data
 - Different properties can be associated with a secondary set,
 - E.g., secondary-only, hidden delayed, arbiters, non-voting
- If the primary fails the secondary sets “vote” to elect the new primary set



Sync vs Async Replication

- **Synchronous replication**: updates are propagated to other replicas as part of a single transaction
- Implementations
 - 2-Phase Commit (2PC)
 - Paxos
 - Both solutions are complex / expensive
- **Asynchronous replication**
 - The primary node propagates updates to replicas
 - The transaction is completed before replicas are updated (even if there are failures)
 - Commits are quick at cost of consistency
- **Eventual consistency**
 - Popularized by AWS DynamoDB
 - Consistency guaranteed only on the eventual outcome
 - “*Eventual*” can mean after the server or network is fixed

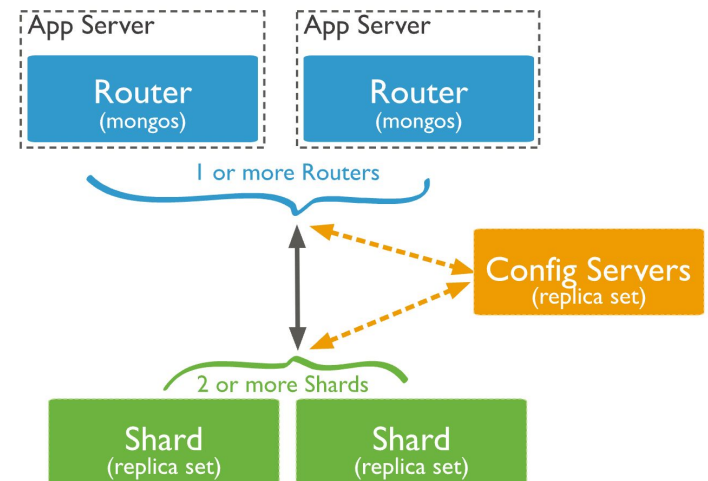
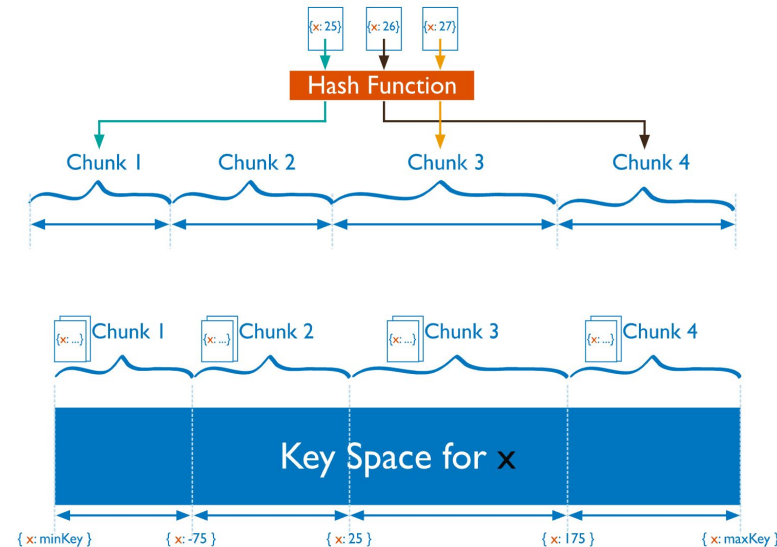


Data Consistency




- Client decides how to enforce consistency for reads
- Reads to a primary have **strict consistency**
 - Reads reflect the latest changes to the data
 - All writes and *consistent* reads go to the primary
- Reads to a secondary have **eventual consistency**
 - Updates propagate gradually
 - Client may read a previous state of the database
 - All *eventually consistent* reads are distributed among the secondaries

MongoDB: Sharding

- **Shard** = subset of data
 - A collection is split in pieces based on the shard key
 - Data distributed based on shard key or intervals [a, b)
- **Sharding** = method for distributing data across different machines
- **Horizontal scaling** can be achieved through sharding
 - Divide data and workload over multiple servers
 - Complexity in infrastructure and maintenance
- **mongos** acts as a query router interfacing clients and sharded cluster
 - Each shard can be deployed as a replica set
 - Config servers store metadata and configuration settings for cluster



RDBMs Internals



Query Processing Engine

- Given a user query, decide how to “execute” it
- Specify sequence of pages to be brought in memory
- Operate upon the tuples to produce results

Buffer Manager




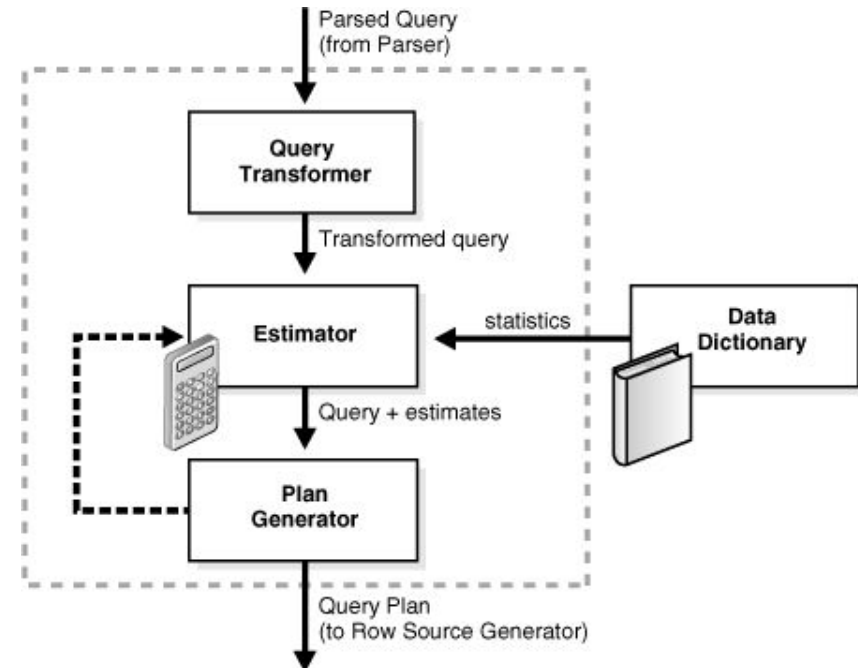
- Bring pages from disk to memory
- Manage the limited memory

Storage hierarchy

- How are tables mapped to files?
- How are tuples mapped to disk blocks?

Query Optimizer

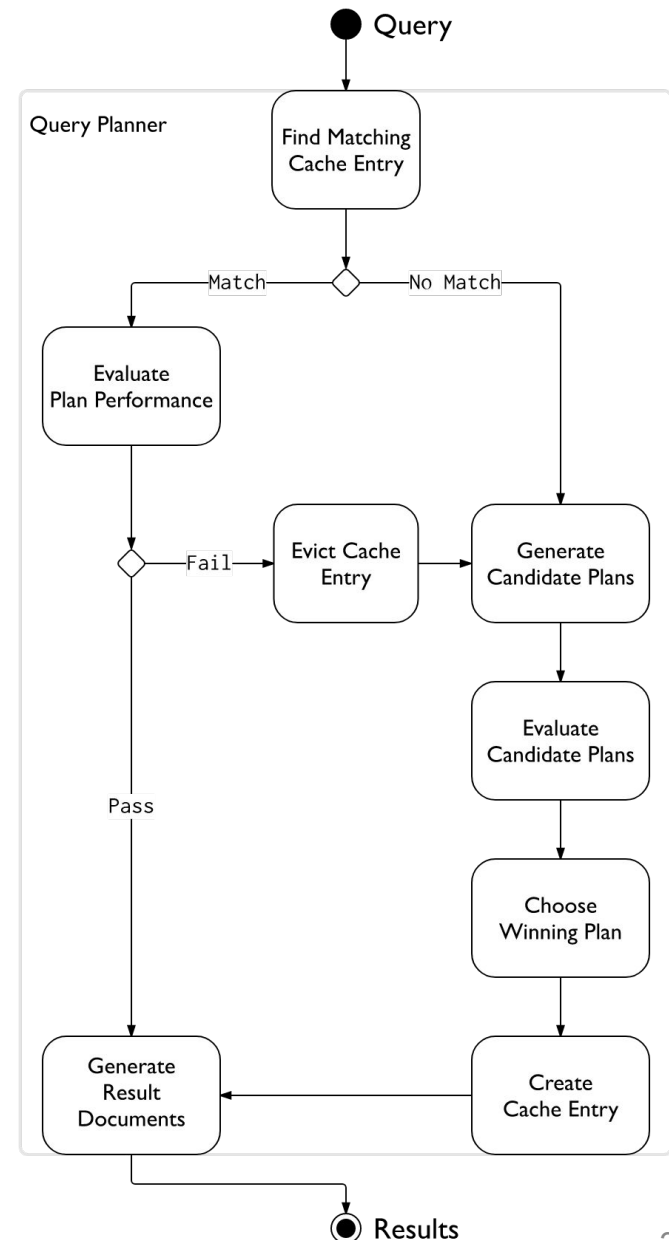
- **RDBMSs: query optimizer is static**
 - Assign a cost to each query plan
 - Estimate some cost params (e.g., time to access data)
 - Search for the best query
 - At least traditional RDBMSs



Query Optimizer

- **MongoDB: query optimizer is dynamic**

- Try different query plans and learn which ones perform well
- The space of query plans is not so large, because there are no joins
- When testing new plans
 - Execute multiple query plans in parallel
 - As soon as one plan finishes, terminate the other plans
- Cache the result
- If a plan that was working well starts performing poorly try again different plans
 - E.g, data in the DB has changed, parameter values to a query are different

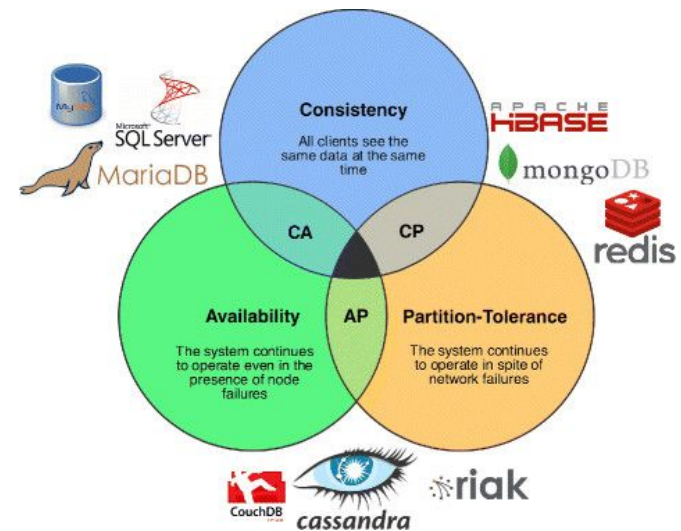


MongoDB: Strengths

- Provide a flexible and modern query language
- High-performance
 - Implemented in C++
- Very rapid development, open source
 - Support for many platforms
 - Many language drivers
- Built to address a distributed database system
 - Sharding
 - Replica sets of data
- Tunable consistency
- Useful for working with a huge quantity of data not requiring a relational model
 - The relationships between the elements does not matter
 - What matters is the ability to store and retrieve great quantities of data

MongoDB: Limitations

- No referential integrity
 - Aka foreign key constraint
- Lack of transactions and joins
- High degree of denormalization
 - Need to update data in many places instead of one
- Lack of predefined schema is a double-edged sword
 - You must have a data model in your application
 - Objects within a collection can be completely inconsistent in their fields
- CAP Theorem: targets consistency and partition tolerance, giving up on availability



UMD DATA605 - Big Data Systems

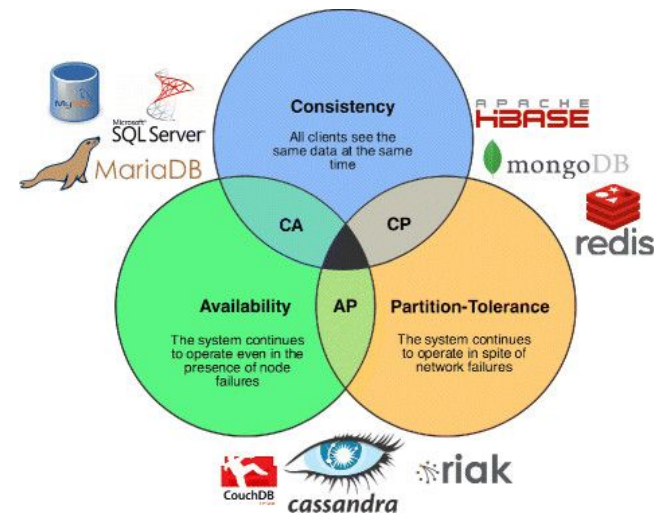
NoSQL document stores

MongoDB

Couchbase

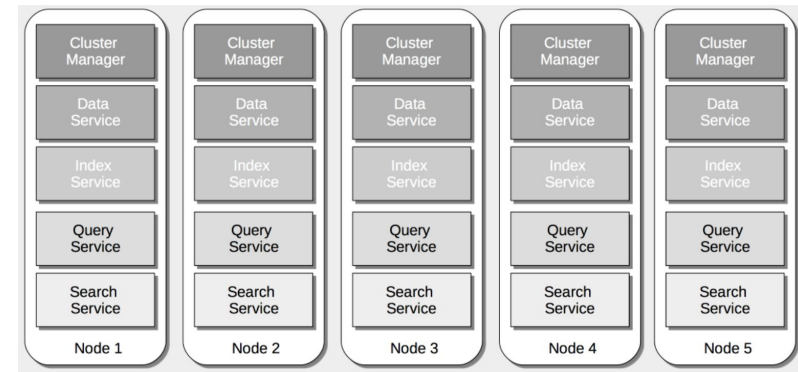

Couchbase

- NoSQL document-oriented DB (like MongoDB)
- [Couchbase](#) = merge of CouchDB and membase
 - *CouchDB*
 - Open source document store
 - HTTP RESTful API to add, update, delete documents
 - Support all 4 ACID properties
 - *membase*
 - Distributed key-value store (like Redis)
 - Designed to scale both up and down
 - Highly available and partition tolerant
 - Uses HTTP protocol to query and interact with objects in the DB
 - No query language
 - Objects stored in *buckets*
 - Collection of JSON docs, with no special relation to one another
- From CAP point of view:
 - Supports consistency and partition tolerance
 - High availability is achieved through use of multiple clusters



Architecture

- Every Couchbase node consists of **different services**:
 - Data service
 - Index service
 - Query service
 - Cluster manager component
- Services can run on separate nodes of the cluster, if needed
- **Data replication**
 - Across nodes of a cluster
 - Across data centers
- **Data service**
 - Writes data *asynchronously* to disk after acknowledging to the client
 - Optionally *synchronous*: ensure data is written to more than one server before acknowledging a write



Queries

- **Can create multiple views over documents**
 - Views are optimized / indexed by Couchbase for fast queries
 - Re-indexed when underlying documents changes
 - Can do full-text searches using the indexes
- **Perform well when:**
 - There are infrequent changes to the structure of documents
 - Know in advance what kinds of queries you want to execute
- **Query**
 - Uses a custom query language called [N1QL](#) (“nickel”)
 - Extends SQL to JSON documents
 - Queries over multiple documents using (server-side) joins
- **Map-reduce support**
 - (Map) First define a view with the columns of the document your are interested in
 - (Reduce) Optionally define aggregate functions over the data

Couchbase vs MongoDB

- According to Couchbase advocates
- **MongoDB**: hard to scale from single replica set to fully distributed environment
- **MongoDB**: require a 3rd party cache to help it perform well
 - **Couchbase**: has integrated in-memory cache (memcached)
 - Keeps frequently accessed documents, metadata, and indexes in RAM, yielding high read/write throughput at low latency
- **MongoDB**: performance degrades with increasing numbers of clients/users
 - **Couchbase**: Scales seamlessly, with an in-memory architecture, and able to scale across multiple nodes
- **MongoDB**: susceptible to data loss from failures
 - **Couchbase**: no master, no single point of failure
 - During failover, prevents different nodes from accepting simultaneous reads of writes of same data (to maintain consistency)